

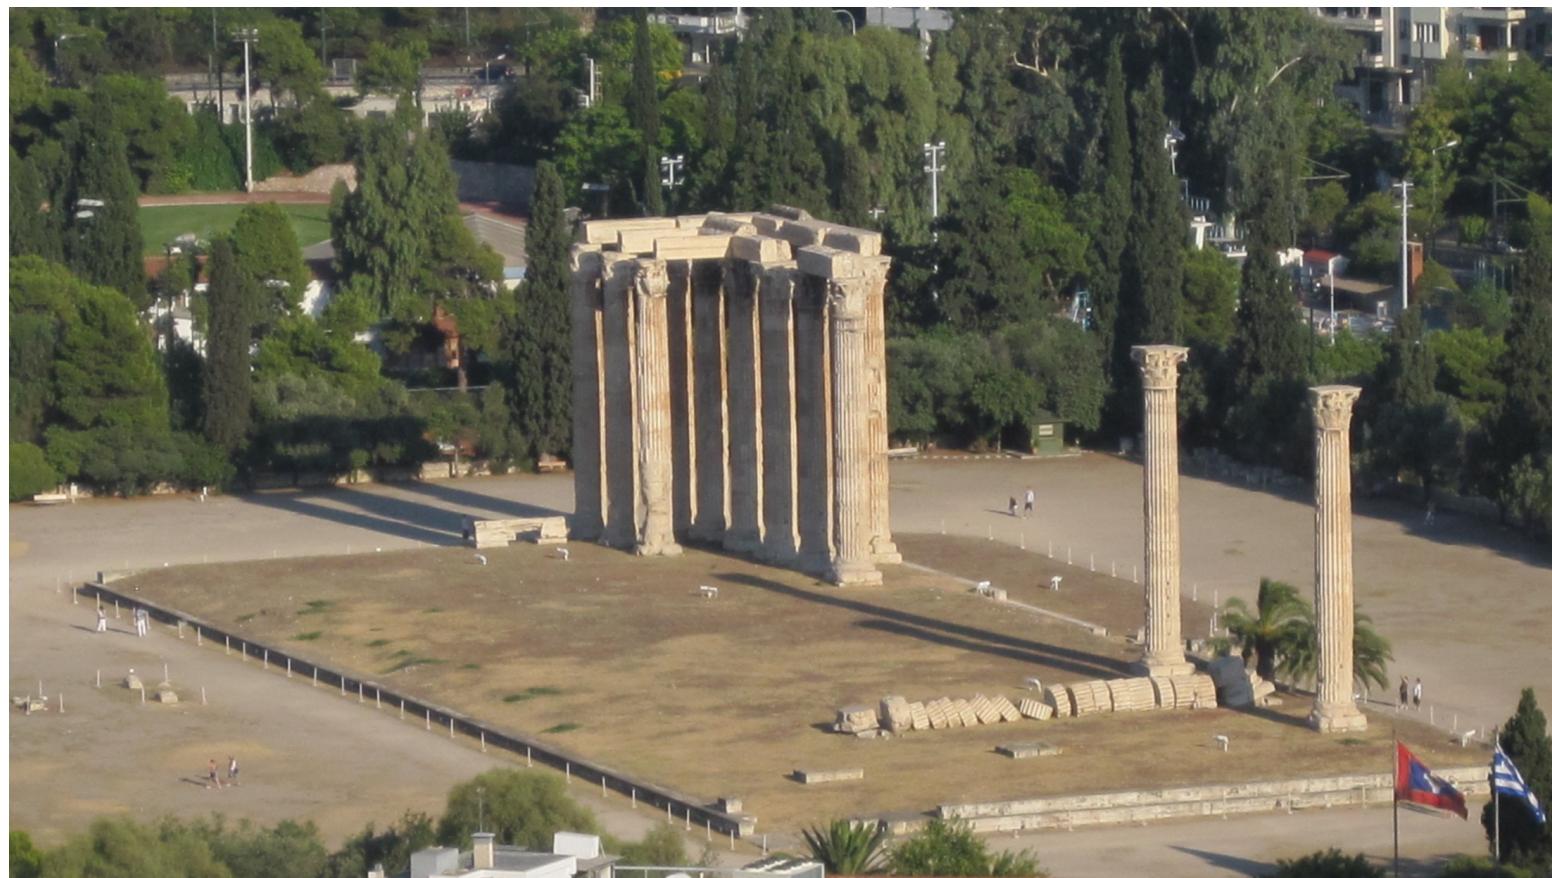


# CS 412 Intro. to Data Mining

## Chapter 4. Data Warehousing and On-line Analytical Processing

(OLAP)

Jiawei Han, Computer Science, Univ. Illinois at Urbana-Champaign, 2017



# Chapter 4: Data Warehousing and On-line Analytical Processing

- Data Warehouse: Basic Concepts
- Data Warehouse Modeling: **Data Cube** and **OLAP**
- Data Warehouse Design and Usage
- Data Warehouse Implementation
- Summary

«#»

## What is a Data Warehouse?

- Defined in many different ways, but not rigorously
  - A decision support database that is maintained **separately** from the organization's operational database
  - Support **information processing** by providing a solid platform of consolidated, historical data for analysis
  - "A data warehouse is a **subject-oriented**, **integrated**, **time-variant**, and **nonvolatile** collection of data in support of management's **decision-making process**."—W. H. Inmon
- Data warehousing:
  - The process of constructing and using data warehouses

«#»

Data ห้องน้ำดูด  
Capture มาจาก  
ก่อสร้าง ไม่ซับซ้อนแล้ว  
อยู่กับที่ statics)

เมื่อ Data Warehouse  
ถูกดึงมาในลักษณะ  
ในครั้ง Data warehouse  
ว่าทำให้เพื่อต้องคำนวณ  
หรือสร้างมาบังเอิญ

Data ห้องน้ำเป็น  
ร่วมเวลาในบังคับ  
อาจไม่สามารถตอบ  
คำถามของคนใดๆ

Data ห้องน้ำ  
ประมวลผล:  
ก่อสร้าง  
เก็บ, กัน  
คงทน  
คงทน

## Data Warehouse—Subject-Oriented

- Organized around major subjects, such as **customer, product, sales**
- Focusing on the modeling and analysis of data for decision makers, not on daily operations or transaction processing
- Provide **a simple and concise** view around particular subject issues by **excluding** data that are not useful in the decision support process

<#>

## Data Warehouse—Integrated

↳ ຈະຕູ້ງຮັບຂໍາມົນຈາກທຳນາລັດໆ ໃຫຍ່ງ

- Constructed by integrating multiple, heterogeneous data sources
  - relational databases, flat files, on-line transaction records
- Data cleaning and data integration techniques are applied.
  - Ensure consistency in naming conventions, encoding structures, attribute measures, etc. among different data sources
  - Ex. Hotel price: differences on currency, tax, breakfast covered, and parking
  - When data is moved to the warehouse, it is converted

<#>

# Data Warehouse—Time Variant

---

- The time horizon for the data warehouse is significantly longer than that of operational systems
  - Operational database: current value data
  - Data warehouse data: provide information from a historical perspective (e.g., past 5-10 years)
- Every key structure in the data warehouse
  - Contains an element of time, explicitly or implicitly
  - But the key of operational data may or may not contain “time element”

‹#›

# Data Warehouse—Nonvolatile

---

- Independence
  - A **physically separate store** of data transformed from the operational environment
- Static: Operational **update of data does not occur** in the data warehouse environment
  - Does not require transaction processing, recovery, and concurrency control mechanisms
  - Requires only two operations in data accessing:
    - *initial loading of data* and *access of data*

‹#›

# OLTP vs. OLAP

- OLTP: Online transactional processing
  - DBMS operations
  - Query and transactional processing
- OLAP: Online analytical processing
  - Data warehouse operations
  - Drilling, slicing, dicing, etc.

	OLTP	OLAP
<b>users</b>	clerk, IT professional	knowledge worker
<b>function</b>	day to day operations	decision support
<b>DB design</b>	application-oriented	subject-oriented
<b>data</b>	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
<b>usage</b>	repetitive	ad-hoc
<b>access</b>	read/write index/hash on prim. key	lots of scans
<b>unit of work</b>	short, simple transaction	complex query
<b># records accessed</b>	tens	millions
<b>#users</b>	thousands	hundreds
<b>DB size</b>	100MB-GB	100GB-TB
<b>metric</b>	transaction throughput	query throughput, response

<#>

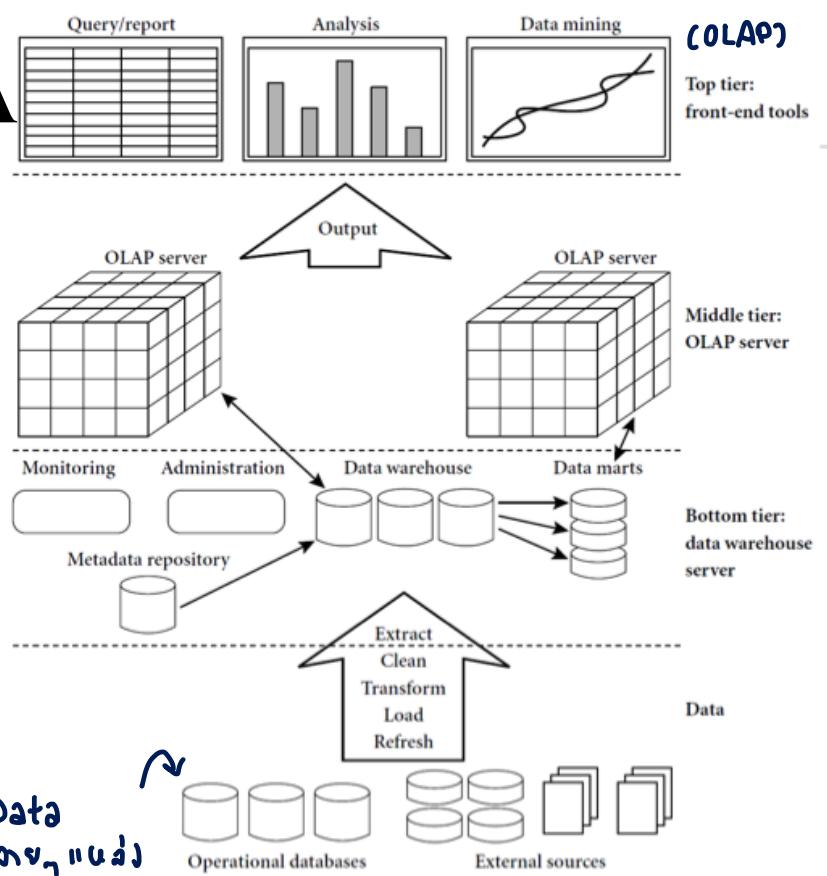
## Why a Separate Data Warehouse?

- High performance for both systems
  - DBMS— tuned for OLTP: access methods, indexing, concurrency control, recovery
  - Warehouse—tuned for OLAP: complex OLAP queries, multidimensional view, consolidation
- Different functions and different data:
  - missing data: Decision support requires historical data which operational DBs do not typically maintain
  - data consolidation: DS requires consolidation (aggregation, summarization) of data from heterogeneous sources
  - data quality: different sources typically use inconsistent data representations, codes and formats which have to be reconciled
- Note: There are more and more systems which perform OLAP analysis directly on *relational databases*

<#>

# Data Warehouse: A Multi-Tiered Architecture

- Top Tier: Front-End Tools
- Middle Tier: OLAP Server
- Bottom Tier: Data Warehouse Server
- Data



## Three Data Warehouse Models

- **Enterprise warehouse**
  - Collects all of the information about subjects spanning the entire organization
- **Data Mart**
  - A subset of corporate-wide data that is of value to a specific group of users
  - Its scope is confined to specific, selected groups, such as marketing data mart
    - Independent vs. dependent (directly from warehouse) data mart
- **Virtual warehouse**
  - A set of views over operational databases
  - Only some of the possible summary views may be materialized

# Extraction, Transformation, and Loading (ETL)

---

- **Data extraction**
  - get data from multiple, heterogeneous, and external sources
- **Data cleaning**
  - detect errors in the data and rectify them when possible
- **Data transformation**
  - convert data from legacy or host format to warehouse format
- **Load**
  - sort, summarize, consolidate, compute views, check integrity, and build indices and partitions
- **Refresh**
  - propagate the updates from the data sources to the warehouse

«#»

## Metadata Repository

---

- **Meta data** is the data defining warehouse objects. It stores:
  - Description of the structure of the data warehouse
    - schema, view, dimensions, hierarchies, derived data defn, data mart locations and contents
  - Operational meta-data
    - data lineage (history of migrated data and transformation path), currency of data (active, archived, or purged), monitoring information (warehouse usage statistics, error reports, audit trails)
    - The algorithms used for summarization
    - The mapping from operational environment to the data warehouse
    - Data related to system performance
      - warehouse schema, view and derived data definitions
    - Business data
      - business terms and definitions, ownership of data, charging policies

«#»

# Chapter 4: Data Warehousing and On-line Analytical Processing

- Data Warehouse: Basic Concepts
- Data Warehouse Modeling: Data Cube and OLAP
- Data Warehouse Design and Usage
- Data Warehouse Implementation
- Summary

<#>

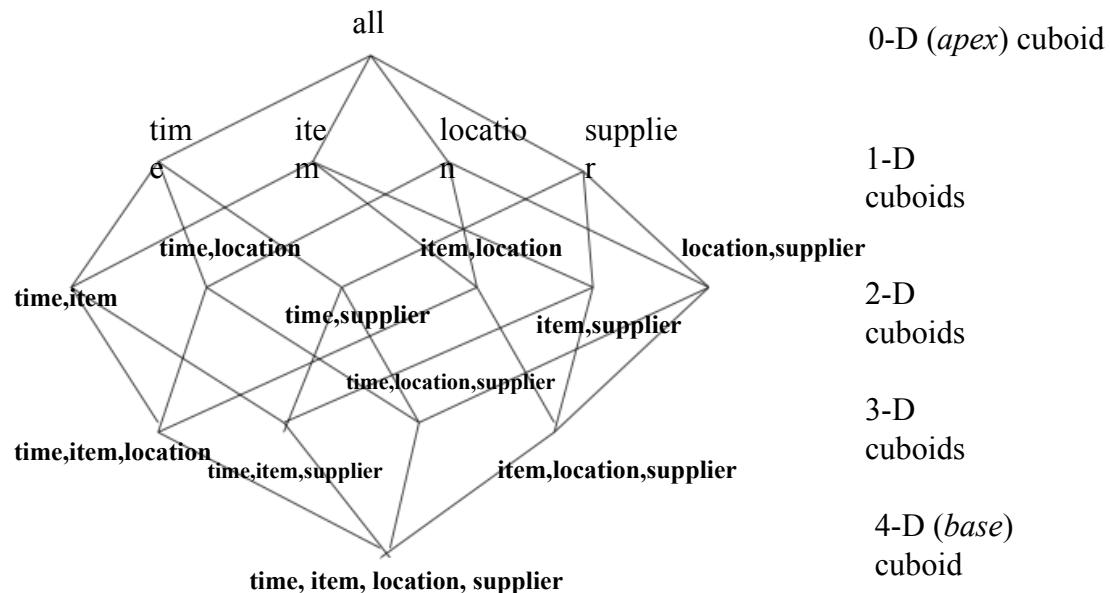
# សំគាល់ #

## From Tables and Spreadsheets to Data Cubes

- A **data warehouse** is based on a multidimensional data model which views data in the form of a data cube
- A data cube, such as sales, allows data to be modeled and viewed in multiple dimensions  
  - **Dimension tables**, such as item (item\_name, brand, type), or time(day, week, month, quarter, year)  
        មានអ្នកផ្តល់ព័ត៌មាន
  - **Fact table** contains **measures** (such as dollars\_sold) and keys to each of the related dimension tables  
        មានរបៀបផ្តល់ព័ត៌មាន
- **Data cube**: A lattice of cuboids
  - In data warehousing literature, an n-D base cube is called a **base cuboid**
  - The top most 0-D cuboid, which holds the highest-level of summarization, is called the **apex cuboid**
  - The lattice of cuboids forms a **data cube**.

<#>

# Data Cube: A Lattice of Cuboids



17

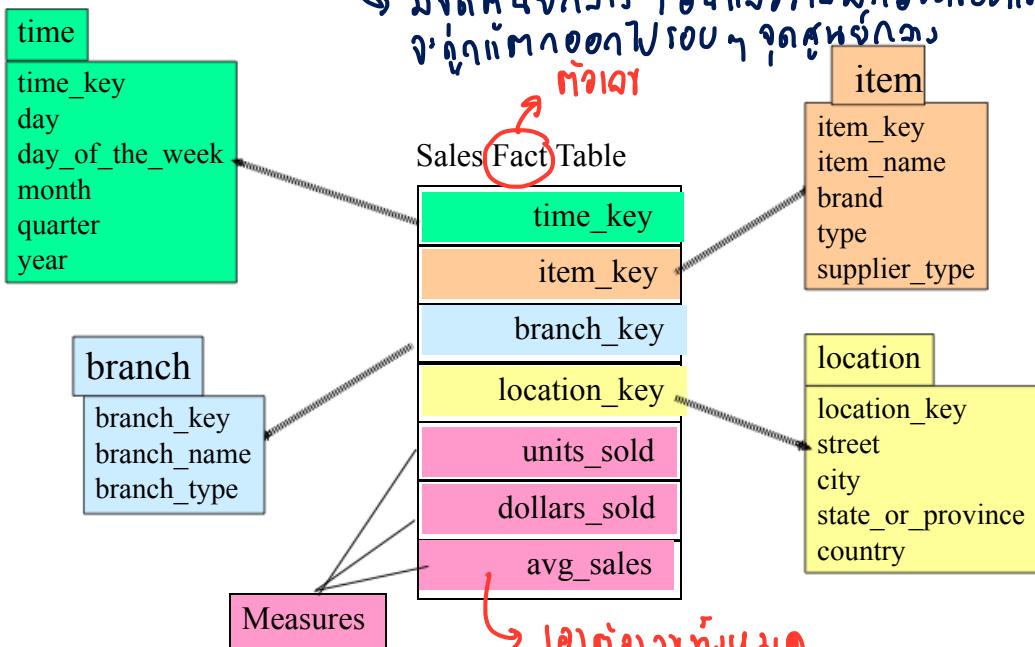
<#>

## Conceptual Modeling of Data Warehouses

- Modeling data warehouses: dimensions & measures
- Star schema: A fact table in the middle connected to a set of dimension tables
- Snowflake schema: A refinement of star schema where some dimensional hierarchy is normalized into a set of smaller dimension tables, forming a shape similar to snowflake
- Fact constellations: Multiple fact tables share dimension tables, viewed as a collection of stars, therefore called **galaxy schema** or fact constellation

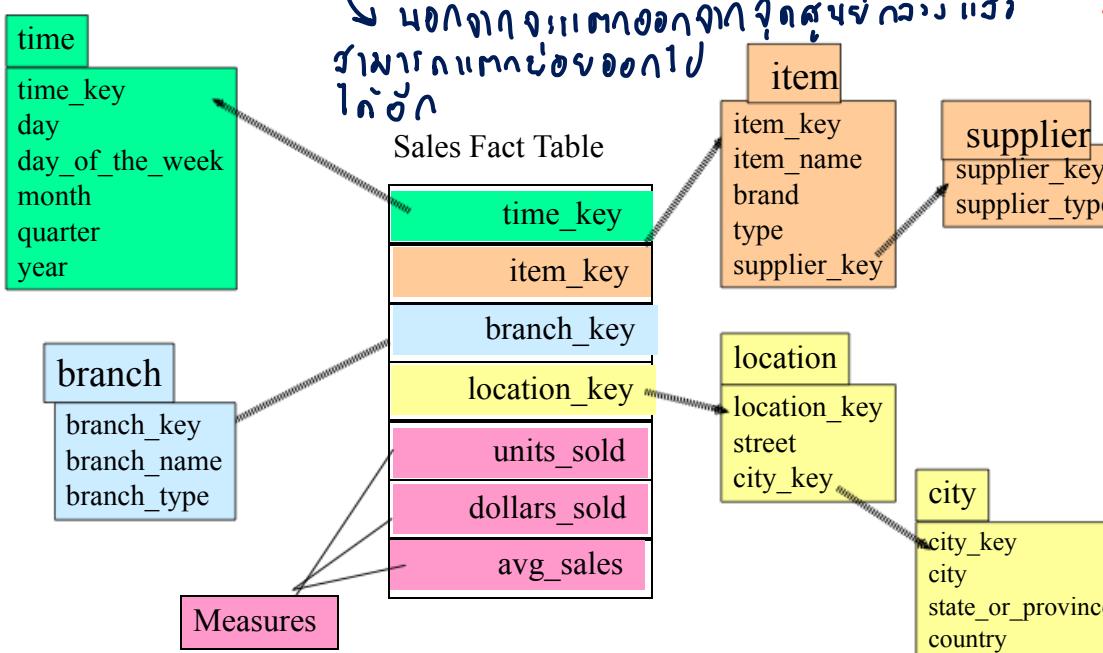
<#>

# Star Schema: An Example



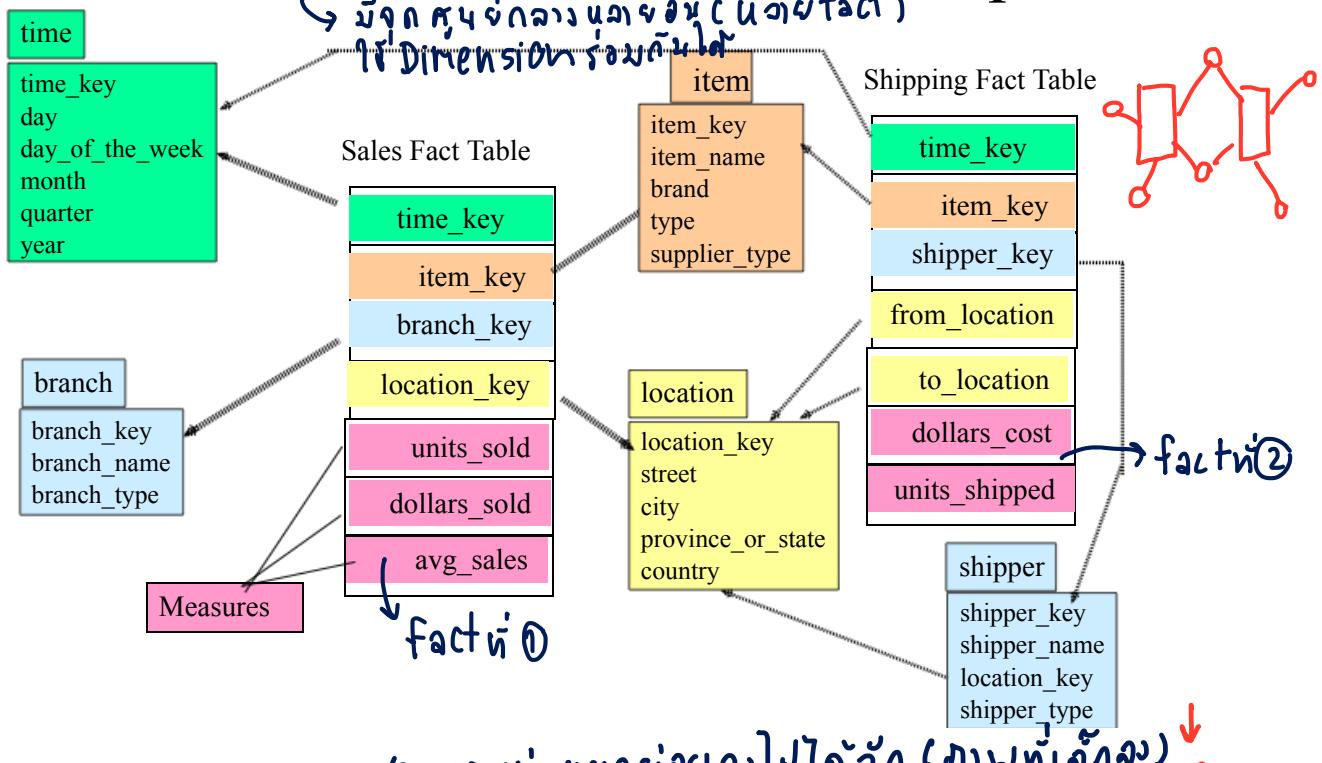
19

# Snowflake Schema: An Example

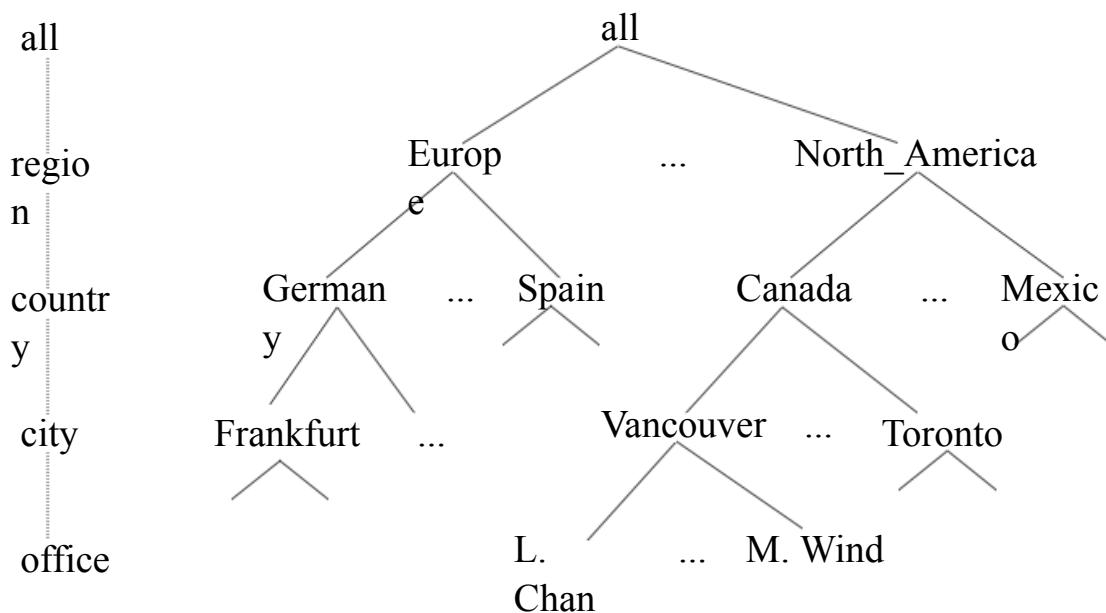


19

# Fact Constellation: An Example



## A Concept Hierarchy for a Dimension (location)

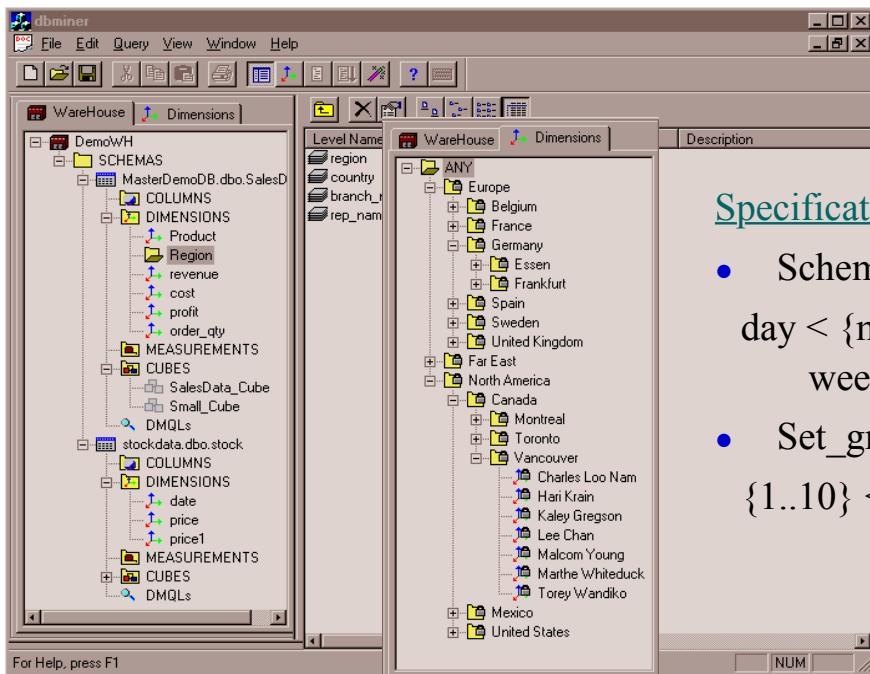


# Data Cube Measures: Three Categories

- Distributive: if the result derived by applying the function to  $n$  aggregate values is the same as that derived by applying the function on all the data without partitioning
  - E.g., count(), sum(), min(), max()
- Algebraic: if it can be computed by an algebraic function with  $M$  arguments (where  $M$  is a bounded integer), each of which is obtained by applying a distributive aggregate function
  - $\text{avg}(x) = \text{sum}(x) / \text{count}(x)$
  - Is  $\text{min\_N}()$  an algebraic measure? How about  $\text{standard\_deviation}()$ ?
- Holistic: if there is no constant bound on the storage size needed to describe a subaggregate.
  - E.g., median(), mode(), rank()

&lt;#&gt;

## View of Warehouses and Hierarchies



### Specification of hierarchies

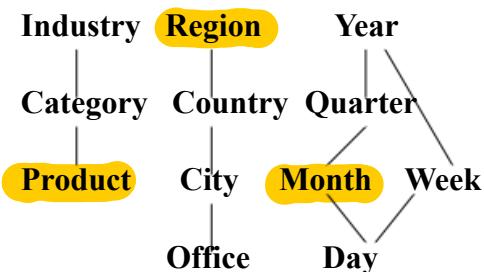
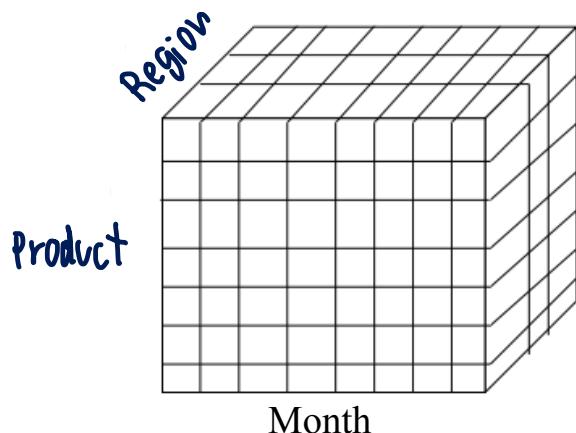
- Schema hierarchy
  - day < {month < quarter; week} < year
- Set\_grouping hierarchy
  - {1..10} < inexpensive

&lt;#&gt;

# Multidimensional Data

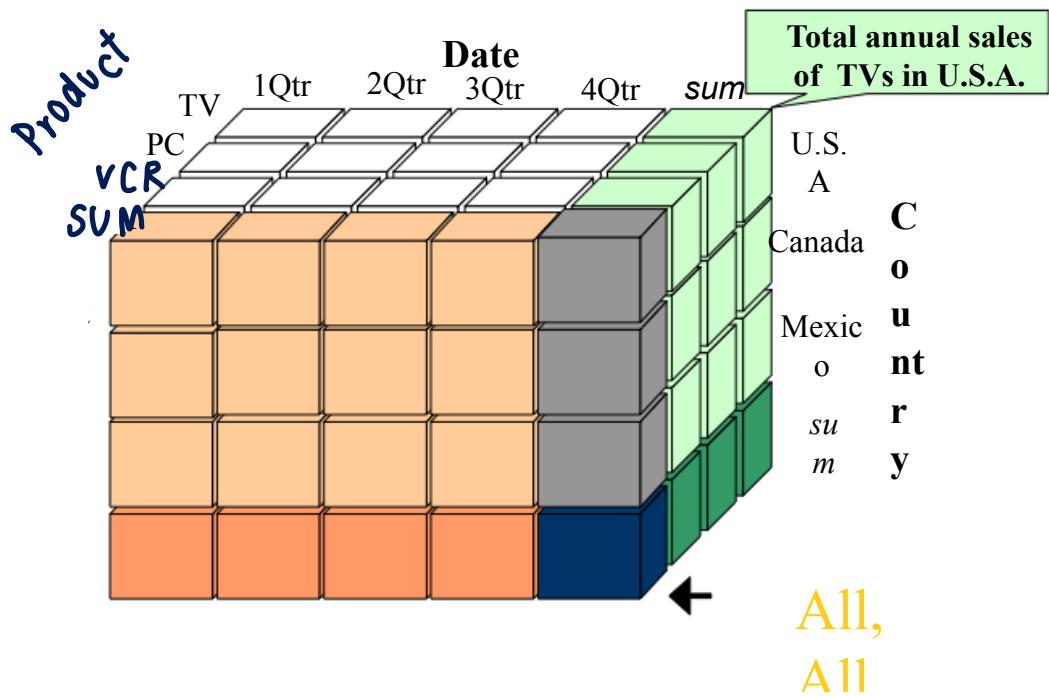
- Sales volume as a function of product, month, and region

Dimensions: *Product, Location, Time*  
Hierarchical summarization paths



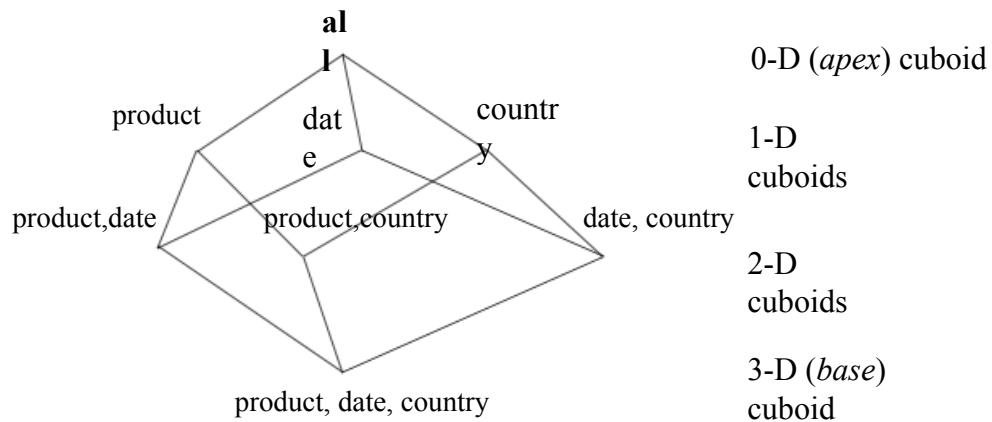
<#>

## A Sample Data Cube



<#>

# Cuboids Corresponding to the Cube

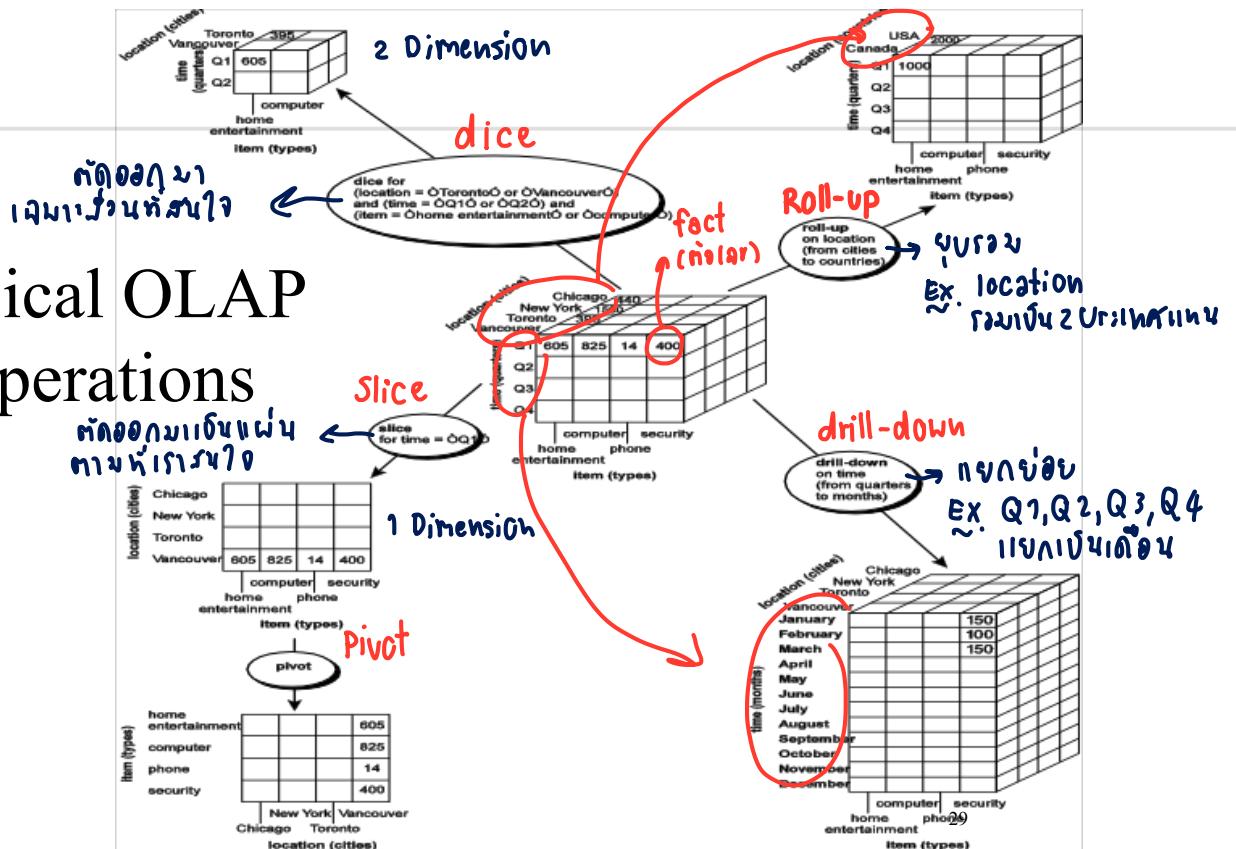


<#>

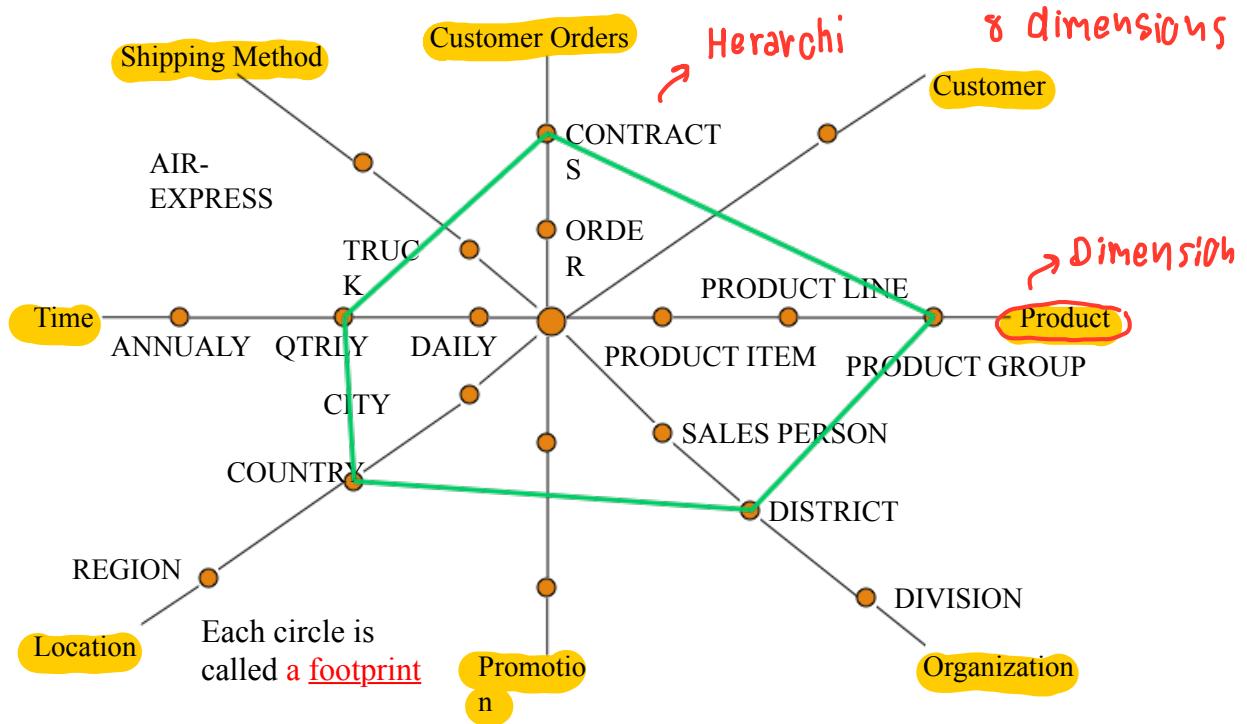
## Typical OLAP Operations

- **Roll up (drill-up):** summarize data
  - *by climbing up hierarchy or by dimension reduction*
- **Drill down (roll down):** reverse of roll-up
  - *from higher level summary to lower level summary or detailed data, or introducing new dimensions*
- **Slice and dice:** *project and select*
- **Pivot (rotate):**
  - *reorient the cube, visualization, 3D to series of 2D planes*
- Other operations
  - **Drill across:** *involving (across) more than one fact table*
  - **Drill through:** *through the bottom level of the cube to its back-end relational tables (using SQL)*

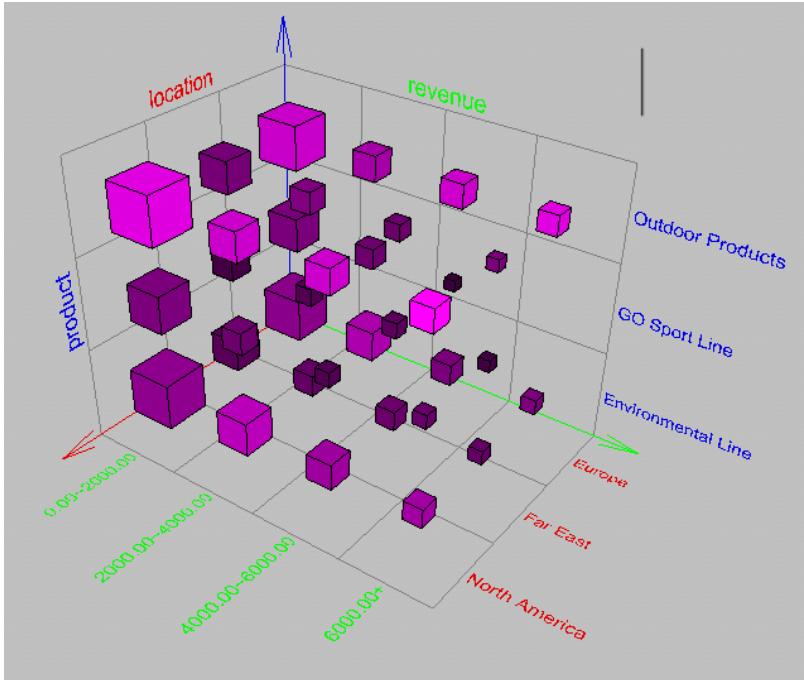
<#>



## A Star-Net Query Model



# Browsing a Data Cube



- Visualization
- OLAP capabilities
- Interactive manipulation

<#>

## Chapter 4: Data Warehousing and On-line Analytical Processing

- Data Warehouse: Basic Concepts
- Data Warehouse Modeling: Data Cube and OLAP
- Data Warehouse Design and Usage
- Data Warehouse Implementation
- Summary

<#>

# Design of Data Warehouse: A Business Analysis Framework

---

- Four views regarding the design of a data warehouse
  - Top-down view
    - allows selection of the relevant information necessary for the data warehouse
  - Data source view
    - exposes the information being captured, stored, and managed by operational systems
  - Data warehouse view
    - consists of fact tables and dimension tables
  - Business query view
    - sees the perspectives of data in the warehouse from the view of end-user

«#»

## Data Warehouse Design Process

---

- **Top-down, bottom-up approaches or a combination of both**
  - Top-down: Starts with overall design and planning (mature)
  - Bottom-up: Starts with experiments and prototypes (rapid)
- **From software engineering point of view**
  - Waterfall: structured and systematic analysis at each step before proceeding to the next
  - Spiral: rapid generation of increasingly functional systems, short turn around time, quick turn around
- **Typical data warehouse design process**
  - Choose a business process to model, e.g., orders, invoices, etc.
  - Choose the *grain* (*atomic level of data*) of the business process
  - Choose the dimensions that will apply to each fact table record
  - Choose the measure that will populate each fact table record

«#»

# Data Warehouse Usage

---

- Three kinds of data warehouse applications
  - **Information processing**
    - supports querying, basic statistical analysis, and reporting using crosstabs, tables, charts and graphs
  - **Analytical processing**
    - multidimensional analysis of data warehouse data
    - supports basic OLAP operations, slice-dice, drilling, pivoting
  - **Data mining**
    - knowledge discovery from hidden patterns
    - supports associations, constructing analytical models, performing classification and prediction, and presenting the mining results using visualization tools

‹#›

## From On-Line Analytical Processing (OLAP) to On Line Analytical Mining (OLAM)

---

- Why **online analytical mining?**
  - High quality of data in data warehouses
  - DW contains integrated, consistent, cleaned data
  - Available information processing structure surrounding data warehouses
  - ODBC, OLEDB, Web accessing, service facilities, reporting and OLAP tools
  - OLAP-based exploratory data analysis
  - Mining with drilling, dicing, pivoting, etc.
  - On-line selection of data mining functions
  - Integration and swapping of multiple mining functions, algorithms, and tasks

‹#›

# Chapter 4: Data Warehousing and On-line Analytical Processing

- Data Warehouse: Basic Concepts
- Data Warehouse Modeling: Data Cube and OLAP
- Data Warehouse Design and Usage
- Data Warehouse Implementation
- Summary

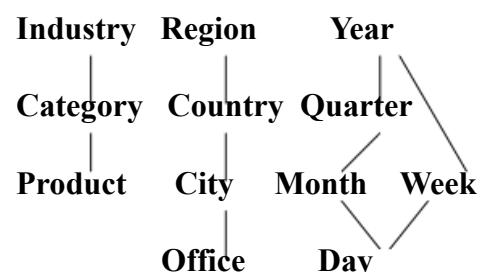
«#»

## Efficient Data Cube Computation

- Data cube can be viewed as a lattice of cuboids
  - The bottom-most cuboid is the base cuboid
  - The top-most cuboid (apex) contains only one cell
  - How many cuboids in an n-dimensional cube with L levels?
- Materialization of data cube
  - **Full materialization:** Materialize every (cuboid)
  - **No materialization:** Materialize none (cuboid)
  - **Partial materialization:** Materialize some cuboids
    - Which cuboids to materialize?
      - Selection based on size, sharing, access frequency, etc.

Why this formula?

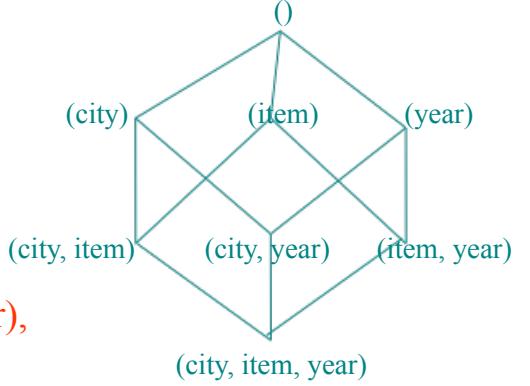
$$T = \prod_{i=1}^n (L_i + 1)$$



«#»

# The “Compute Cube” Operator

- Cube definition and computation in DMQL  
`define cube sales [item, city, year]: sum (sales_in_dollars)`  
`compute cube sales`
- Transform it into a SQL-like language (with a new operator `cube by`, introduced by Gray et al.’96)  
SELECT item, city, year, SUM (amount)  
FROM SALES  
`CUBE BY` item, city, year
- Need compute the following Group-Bys  
`(date, product, customer),`  
`(date, product),(date, customer), (product, customer),`  
`(date), (product), (customer)`



«#»

◊

## Indexing OLAP Data: Bitmap Index

- Index on a particular column
  - Each value in the column has a bit vector: bit-op is fast
  - The length of the bit vector: # of records in the base table
  - The  $i$ -th bit is set if the  $i$ -th row of the base table has the value for the indexed column
  - not suitable for high cardinality domains
- A recent bit compression technique, Word-Aligned Hybrid (WAH), makes it work for high cardinality domain as well [Wu, et al. TODS’06]

Base table

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

Index on

RecID	Asia	Europe	America
1	1	0	0
2	0	1	0
3	1	0	0
4	0	0	1
5	0	1	0

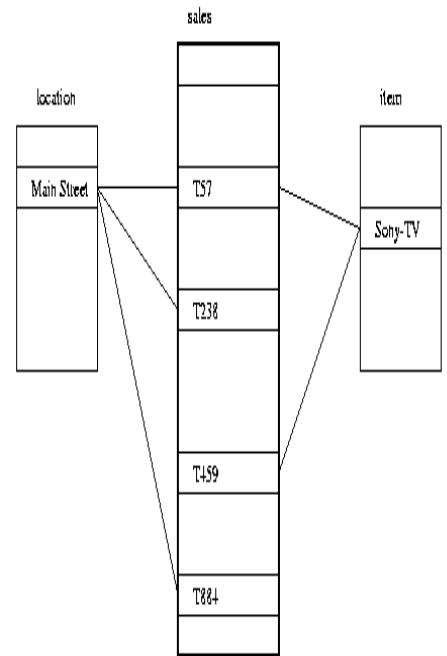
Index on Type

RecID	Retail	Dealer
1	1	0
2	0	1
3	0	1
4	1	0
5	0	1

«#»

# Indexing OLAP Data: Join Indices

- Join index:  $JI(R\text{-id}, S\text{-id})$  where  $R$  ( $R\text{-id}, \dots$ )  $\square \square$   $S$  ( $S\text{-id}, \dots$ )
- Traditional indices map the values to a list of record ids
  - It materializes relational join in JI file and speeds up relational join
- In data warehouses, join index relates the values of the dimensions of a star schema to rows in the fact table.
  - E.g., fact table: *Sales* and two dimensions *city* and *product*
  - A join index on *city* maintains for each distinct city a list of R-IDs of the tuples recording the Sales in the city
- Join indices can span multiple dimensions



## Efficient Processing OLAP Queries

- **Determine which operations** should be performed on the available cuboids
  - Transform drill, roll, etc. into corresponding SQL and/or OLAP operations, e.g., dice = selection + projection
- **Determine which materialized cuboid(s)** should be selected for OLAP op.
  - Let the query to be processed be on  $\{brand, province\_or\_state\}$  with the condition “ $year = 2004$ ”, and there are 4 materialized cuboids available:
    - 1)  $\{year, item\_name, city\}$
    - 2)  $\{year, brand, country\}$
    - 3)  $\{year, brand, province\_or\_state\}$
    - 4)  $\{item\_name, province\_or\_state\}$  where  $year = 2004$Which should be selected to process the query?
- Explore indexing structures and compressed vs. dense array structs in MOLAP

# OLAP Server Architectures

---

- **Relational OLAP (ROLAP)**
    - Use relational or extended-relational DBMS to store and manage warehouse data and OLAP middle ware
    - Include optimization of DBMS backend, implementation of aggregation navigation logic, and additional tools and services
    - Greater scalability
  - **Multidimensional OLAP (MOLAP)**
    - Sparse array-based multidimensional storage engine
    - Fast indexing to pre-computed summarized data
  - **Hybrid OLAP (HOLAP)** (e.g., Microsoft SQLServer)
    - Flexibility, e.g., low level: relational, high-level: array
    - Specialized SQL servers (e.g., Redbricks)
- «#»
- Specialized support for SQL queries over star/snowflake schemas

## Chapter 4: Data Warehousing and On-line Analytical Processing

---

- Data Warehouse: Basic Concepts
- Data Warehouse Modeling: Data Cube and OLAP
- Data Warehouse Design and Usage
- Data Warehouse Implementation
- Summary

# Summary

---

- Data warehousing: A multi-dimensional model of a data warehouse
  - A data cube consists of *dimensions & measures*
  - Star schema, snowflake schema, fact constellations
  - OLAP operations: drilling, rolling, slicing, dicing and pivoting
- Data Warehouse Architecture, Design, and Usage
  - Multi-tiered architecture
  - Business analysis design framework
  - Information processing, analytical processing, data mining, OLAM
- Implementation: Efficient computation of data cubes
  - Partial vs. full vs. no materialization
  - Indexing OALP data: Bitmap index and join index
  - OLAP query processing
  - OLAP servers: ROLAP, MOLAP, HOLAP

«#»

## References (I)

---

- S. Agarwal, R. Agrawal, P. M. Deshpande, A. Gupta, J. F. Naughton, R. Ramakrishnan, and S. Sarawagi. On the computation of multidimensional aggregates. VLDB'96
- D. Agrawal, A. E. Abbadi, A. Singh, and T. Yurek. Efficient view maintenance in data warehouses. SIGMOD'97
- R. Agrawal, A. Gupta, and S. Sarawagi. Modeling multidimensional databases. ICDE'97
- **S. Chaudhuri and U. Dayal. An overview of data warehousing and OLAP technology. ACM SIGMOD Record, 26:65-74, 1997**
- J. Gray, et al. Data cube: A relational aggregation operator generalizing group-by, cross-tab and sub-totals. Data Mining and Knowledge Discovery, 1:29-54, 1997.
- A. Gupta and I. S. Mumick. Materialized Views: Techniques, Implementations, and Applications. MIT Press, 1999
- J. Han. Towards on-line analytical mining in large databases. *SIGMOD Record*, 1998

«#»

## References (II)

- C. Imhoff, N. Galembo, and J. G. Geiger. Mastering Data Warehouse Design: Relational and Dimensional Techniques. John Wiley, 2003
- W. H. Inmon. Building the Data Warehouse. John Wiley, 1996
- R. Kimball and M. Ross. The Data Warehouse Toolkit: The Complete Guide to Dimensional Modeling. 2ed. John Wiley, 2002
- P. O'Neil and D. Quass. Improved query performance with variant indexes. SIGMOD'97
- S. Sarawagi and M. Stonebraker. Efficient organization of large multidimensional arrays. ICDE'94
- P. Valduriez. Join indices. ACM Trans. Database Systems, 12:218-246, 1987.
- J. Widom. Research problems in data warehousing. CIKM'95.
- K. Wu, E. Otoo, and A. Shoshani, Optimal Bitmap Indices with Efficient Compression, ACM Trans. on Database Systems (TODS), 31(1), 2006, pp. 1-38.

«#»

